

1. Complete R Lab 3 from Section 10.6 of our textbook and submit associated R codes and outputs.

Answer->

Code-

```
library (ISLR)
```

```
nci.labs= NCI60$labs
```

```
nci.data= NCI60$data
```

```
dim(nci.data)
```

```
> dim(nci.data)
```

```
[1] 64 6830
```

```
nci.labs [1:4]
```

```
table(nci.labs)
```

```
> nci.labs [1:4]
```

```
[1] "CNS" "CNS" "CNS" "RENAL"
```

```
> table(nci.labs)
```

```
nci.labs
```

```

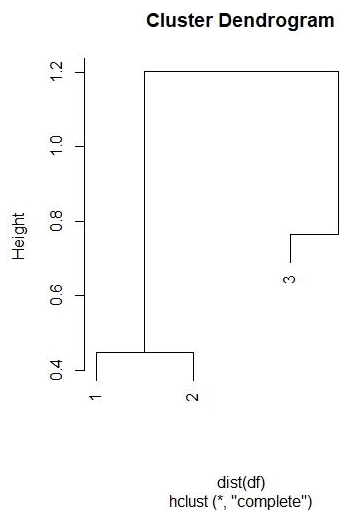
      BREAST      CNS      COLON K562A-repro
        7        5        7        1
K562B-repro LEUKEMIA MCF7A-repro MCF7D-repro
        1        6        1        1
      MELANOMA      NSCLC      OVARIAN      PROSTATE
        8        9        6        2
      RENAL      UNKNOWN
        9        1
```

```
> pr.out =prcomp (nci.data , scale=TRUE)
```

```
> Cols=function (vec ){
```

```
+ + cols=rainbow (length (unique (vec )))
```

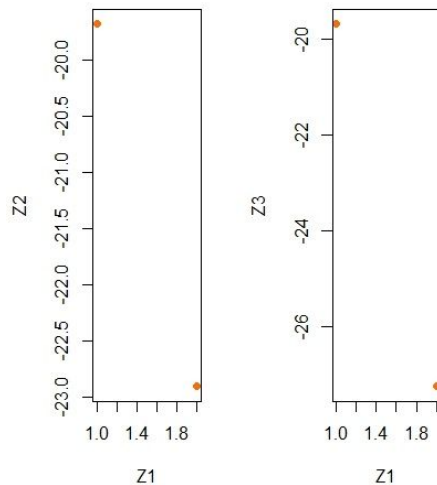
```
+ + return (cols[as.numeric (as.factor (vec))])}
```



```

> Cols=function(vec){cols=rainbow (length(unique(vec)))
+ +   return(cols[as.numeric(as.factor(vec))])}
> par(mfrow =c(1,2))
> plot(pr.out$x [1:2], col=Cols(nci.labs), pch =19,xlab ="Z1",ylab="Z2")
> plot(pr.out$x[c(1,3)], col=Cols(nci.labs), pch =19,xlab ="Z1",ylab="Z3")

```



```

> summary (pr.out)

```

Importance of components:

	PC1	PC2	PC3
Standard deviation	27.8535	21.48136	19.82046
Proportion of Variance	0.1136	0.06756	0.05752
Cumulative Proportion	0.1136	0.18115	0.23867

	PC4	PC5	PC6
Standard deviation	17.03256	15.97181	15.72108
Proportion of Variance	0.04248	0.03735	0.03619
Cumulative Proportion	0.28115	0.31850	0.35468

	PC7	PC8	PC9
Standard deviation	14.47145	13.54427	13.14400
Proportion of Variance	0.03066	0.02686	0.02529
Cumulative Proportion	0.38534	0.41220	0.43750

	PC10	PC11	PC12
Standard deviation	12.73860	12.68672	12.15769
Proportion of Variance	0.02376	0.02357	0.02164
Cumulative Proportion	0.46126	0.48482	0.50646

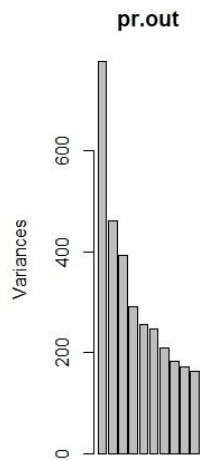
	PC13	PC14	PC15
Standard deviation	11.83019	11.62554	11.43779
Proportion of Variance	0.02049	0.01979	0.01915
Cumulative Proportion	0.52695	0.54674	0.56590

	PC16	PC17	PC18
Standard deviation	11.00051	10.65666	10.48880
Proportion of Variance	0.01772	0.01663	0.01611
Cumulative Proportion	0.58361	0.60024	0.61635

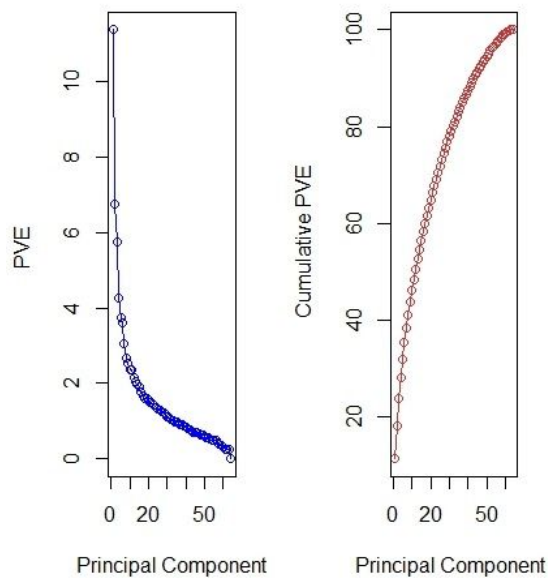
	PC19	PC20	PC21	PC22
Standard deviation				
Proportion of Variance				
Cumulative Proportion				

Standard deviation	10.43518	10.3219	10.14608	10.0544
Proportion of Variance	0.01594	0.0156	0.01507	0.0148
Cumulative Proportion	0.63229	0.6479	0.66296	0.6778
	PC23	PC24	PC25	PC26
Standard deviation	9.90265	9.64766	9.50764	9.33253
Proportion of Variance	0.01436	0.01363	0.01324	0.01275
Cumulative Proportion	0.69212	0.70575	0.71899	0.73174
	PC27	PC28	PC29	PC30
Standard deviation	9.27320	9.0900	8.98117	8.75003
Proportion of Variance	0.01259	0.0121	0.01181	0.01121
Cumulative Proportion	0.74433	0.7564	0.76824	0.77945
	PC31	PC32	PC33	PC34
Standard deviation	8.59962	8.44738	8.37305	8.21579
Proportion of Variance	0.01083	0.01045	0.01026	0.00988
Cumulative Proportion	0.79027	0.80072	0.81099	0.82087
	PC35	PC36	PC37	PC38
Standard deviation	8.15731	7.97465	7.90446	7.82127
Proportion of Variance	0.00974	0.00931	0.00915	0.00896
Cumulative Proportion	0.83061	0.83992	0.84907	0.85803
	PC39	PC40	PC41	PC42
Standard deviation	7.72156	7.58603	7.45619	7.3444
Proportion of Variance	0.00873	0.00843	0.00814	0.0079
Cumulative Proportion	0.86676	0.87518	0.88332	0.8912
	PC43	PC44	PC45	PC46
Standard deviation	7.10449	7.0131	6.95839	6.8663
Proportion of Variance	0.00739	0.0072	0.00709	0.0069
Cumulative Proportion	0.89861	0.9058	0.91290	0.9198
	PC47	PC48	PC49	PC50
Standard deviation	6.80744	6.64763	6.61607	6.40793
Proportion of Variance	0.00678	0.00647	0.00641	0.00601
Cumulative Proportion	0.92659	0.93306	0.93947	0.94548
	PC51	PC52	PC53	PC54
Standard deviation	6.21984	6.20326	6.06706	5.91805
Proportion of Variance	0.00566	0.00563	0.00539	0.00513
Cumulative Proportion	0.95114	0.95678	0.96216	0.96729
	PC55	PC56	PC57	PC58
Standard deviation	5.91233	5.73539	5.47261	5.2921
Proportion of Variance	0.00512	0.00482	0.00438	0.0041
Cumulative Proportion	0.97241	0.97723	0.98161	0.9857
	PC59	PC60	PC61	PC62
Standard deviation	5.02117	4.68398	4.17567	4.08212
Proportion of Variance	0.00369	0.00321	0.00255	0.00244
Cumulative Proportion	0.98940	0.99262	0.99517	0.99761
	PC63	PC64		
Standard deviation	4.04124	2.148e-14		
Proportion of Variance	0.00239	0.000e+00		
Cumulative Proportion	1.00000	1.000e+00		

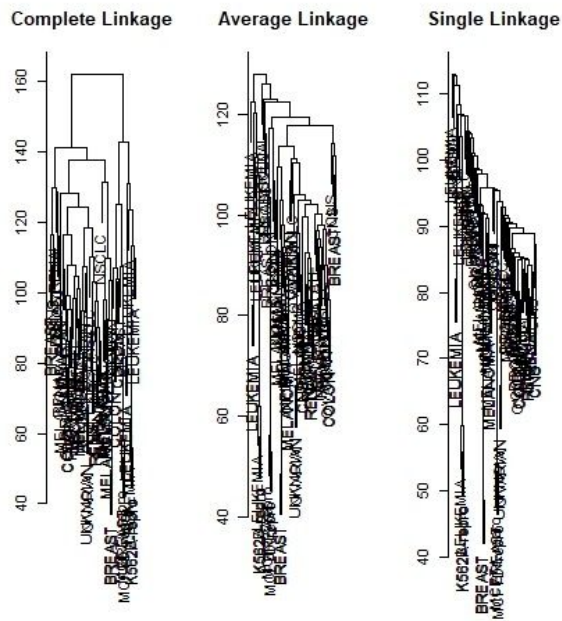
```
> plot(pr.out)
```



```
> pve = 100 * pr.out$sdev ^2 / sum(pr.out$sdev ^2)
> par(mfrow = c(1,2))
> plot(pve , type ="o", ylab="PVE ", xlab=" Principal Component ", col =" blue")
> plot(cumsum (pve ), type="o", ylab =" Cumulative PVE", xlab="Principal Component ", col ="
brown3 ")
```



```
> par(mfrow = c(1,3))
> data.dist=dist(sd.data)
> plot(hclust (data.dist), labels =nci.labs , main=" Complete Linkage ", xlab ="" , sub ="" , ylab = "")
> plot(hclust (data.dist , method ="average"), labels =nci.labs , main=" Average Linkage ", xlab
="" , sub ="" , ylab = "")
> plot(hclust (data.dist , method ="single"), labels =nci.labs , main=" Single Linkage ", xlab="",
sub ="" , ylab = "")
```



```
> hc.out =hclust (dist(sd.data))
> hc.clusters =cutree (hc.out ,4)
> table(hc.clusters ,nci.labs)
```

nci.labs

hc.clusters BREAST CNS COLON K562A-repro K562B-repro

1	2	3	2	0	0
2	3	2	0	0	0
3	0	0	0	1	1
4	2	0	5	0	0

nci.labs

hc.clusters LEUKEMIA MCF7A-repro MCF7D-repro MELANOMA

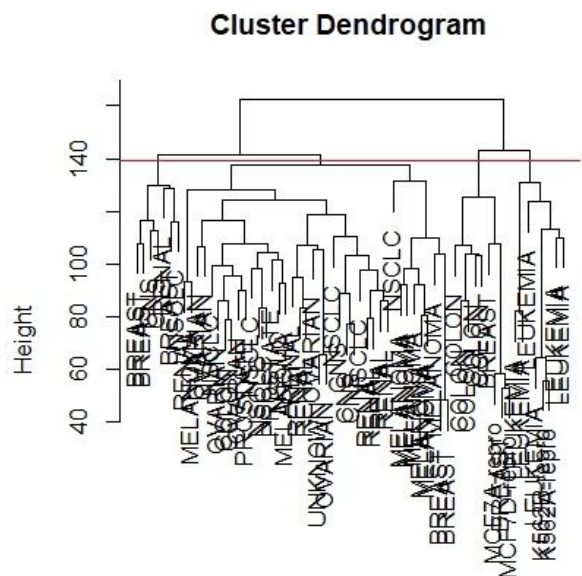
1	0	0	0	8
2	0	0	0	0
3	6	0	0	0
4	0	1	1	0

nci.labs

hc.clusters NSCLC OVARIAN PROSTATE RENAL UNKNOWN

1	8	6	2	8	1
2	1	0	0	1	0
3	0	0	0	0	0
4	0	0	0	0	0

```
> par(mfrow =c(1,1))
> plot(hc.out , labels =nci.labs)
> abline (h=139, col =" red ")
```



```
dist(sd.data)
hclust (*, "complete")
```

```
> hc.out
```

Call:

```
hclust(d = dist(sd.data))
```

Cluster method : complete

Distance : euclidean

Number of objects: 64

```
> set.seed (2)
```

```
> km.out =kmeans (sd.data , 4, nstart =20)
```

```
> km.clusters =km.out$cluster
```

```
> table(km.clusters ,hc.clusters)
```

	hc.clusters			
km.clusters	1	2	3	4
1	11	0	0	9
2	20	7	0	0
3	9	0	0	0
4	0	0	8	0

```
> hc.out =hclust (dist(pr.out$x [ ,1:5]) )
```

```
> plot(hc.out , labels =nci.labs , main=" Hier. Clust . on First Five Score Vectors ")
```

```
> table(cutree (hc.out,4) , nci.labs)
```

	nci.labs					
	BREAST	CNS	COLON	K562A-repro	K562B-repro	LEUKEMIA
1	0	2	7	0	0	2
2	5	3	0	0	0	0
3	0	0	0	1	1	4

4	2	0	0	0	0	0
nci.labs						
MCF7A-repro MCF7D-repro MELANOMA NSCLC OVARIAN						
1	0	0	1	8	5	
2	0	0	7	1	1	
3	0	0	0	0	0	
4	1	1	0	0	0	
nci.labs						
PROSTATE RENAL UNKNOWN						
1	2	7	0			
2	0	2	1			
3	0	0	0			
4	0	0	0			

2. Exercises 2, 3, 9, and 10 from Section 10.7 of our textbook.

2. Suppose that we have four observations, for which we compute a dissimilarity matrix, given by

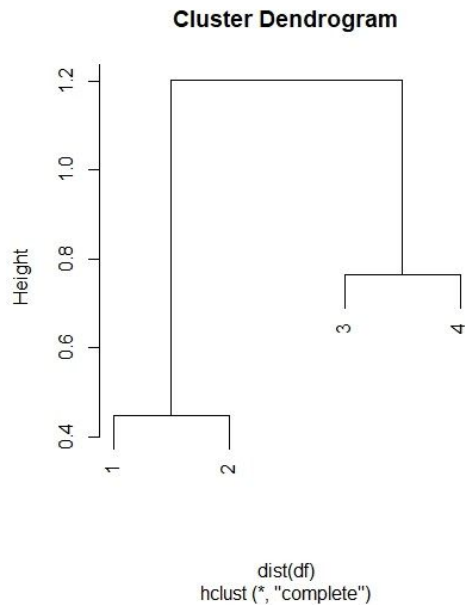
```
[ 0.3 0.4 0.7
 0.3 0.5 0.8
 0.4 0.5 0.45
 0.7 0.8 0.45 ]
```

For instance, the dissimilarity between the first and second observations is 0.3, and the dissimilarity between the second and fourth observations is 0.8.

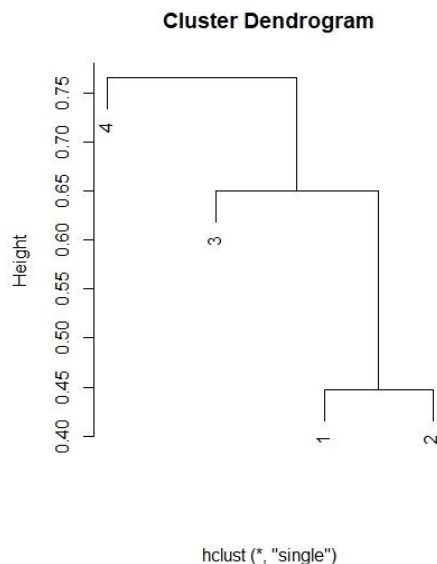
(a) On the basis of this dissimilarity matrix, sketch the dendrogram that results from hierarchically clustering these four observations using complete linkage. Be sure to indicate on the plot the height at which each fusion occurs, as well as the observations corresponding to each leaf in the dendrogram.

Answer->

```
> library(knitr)
> df=data.frame(c(0,0.3,0.4,0.7),c(0.3,0.0,0.5,0.8),c(0.4,0.5,0.0,0.45),c(0.7,0.8,0.45,0.0))
> #print(d)
> colnames(df)=c(paste('Col',1:4))
> kable(df)
| Col 1| Col 2| Col 3| Col 4|
|----:|----:|----:|----:|
| 0.0| 0.3| 0.40| 0.70|
| 0.3| 0.0| 0.50| 0.80|
| 0.4| 0.5| 0.00| 0.45|
| 0.7| 0.8| 0.45| 0.00|
> hcl<- hclust(dist(df))
> plot(hcl)
```



(b) Repeat (a), this time using single linkage clustering.



(c) Suppose that we cut the dendrogram obtained in (a) such that two clusters result. Which observations are in each cluster?

Answer-> If we cut the dendrogram in two clusters, We will obtain (1,2) in first and (3,4) in second.

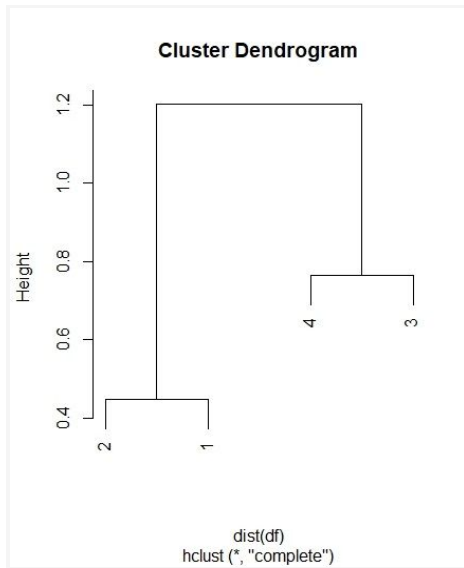
(d) Suppose that we cut the dendrogram obtained in (b) such that two clusters result. Which observations are in each cluster?

Answer-> The clusters obtained here are (4) and (1,2,3).

(e) It is mentioned in the chapter that at each fusion in the dendrogram, the position of the two clusters being fused can be swapped without changing the meaning of the dendrogram. Draw a dendrogram that is equivalent to the dendrogram in (a), for which two or more of the leaves are repositioned, but for which the meaning of the dendrogram is the same.

Answer-> A dendrogram is read bottom up, where the height indicates where clusters are fused. Thus there is no horizontal meaning, the leaves are be swapped but they still represent clusters that are fused at the same height.

```
row.names(DissMatrix)=c(2,1,4,3)
plot(hclust(dist(DissMatrix)))
```



3. In this problem, you will perform K-means clustering manually, with $K = 2$, on a small example with $n = 6$ observations and $p = 2$ features. The observations are as follows.

Obs. X1 X2

1 1 4

2 1 3

3 0 4

4 5 1

5 6 2

6 4 0

(a) Plot the observations.

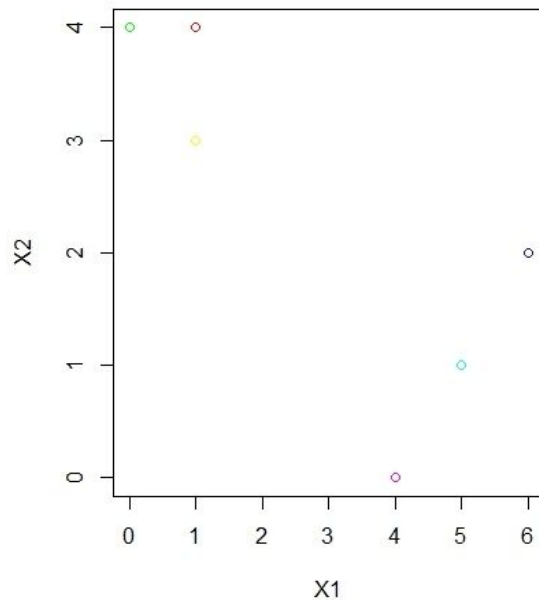
Answer->

```
tb=data.frame(c(1,1,0,5,6,4),c(4,3,4,1,2,0))
```

```
colnames(tb)=c('X1','X2')
```

```
rownames(tb)=1:6
```

```
plot(tb,col=rainbow(6))
```



(b) Randomly assign a cluster label to each observation. You can use the `sample()` command in R to do this. Report the cluster labels for each observation.

Answer->

```
set.seed(1)
> x = cbind(c(1, 1, 0, 5, 6, 4), c(4, 3, 4, 1, 2, 0))
> labels = sample(2, nrow(x), replace=T)
> labels
[1] 1 2 1 1 2 1
```

(c) Compute the centroid for each cluster.

Answer->

```
> centroid1 = c(mean(x[labels==1, 1]), mean(x[labels==1, 2]))
> centroid2 = c(mean(x[labels==2, 1]), mean(x[labels==2, 2]))
> centroid1
[1] 2.50 2.25
> centroid2
[1] 3.5 2.5
```

(d) Assign each observation to the centroid to which it is closest, in terms of Euclidean distance. Report the cluster labels for each Observation.

Answer->

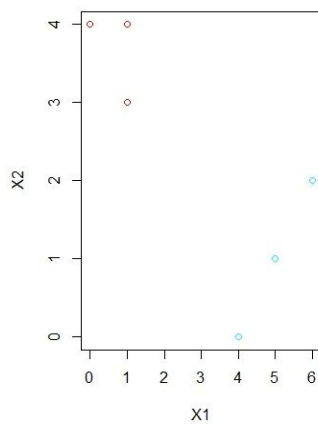
```
> euclid = function(a, b) {
+   return(sqrt((a[1] - b[1])^2 + (a[2]-b[2])^2))
+ }
> assign_labels = function(x, centroid1, centroid2) {
+   labels = rep(NA, nrow(x))
+   for (i in 1:nrow(x)) {
```

```

+   if (euclid(x[i,], centroid1) < euclid(x[i,], centroid2)) {
+     labels[i] = 1
+   } else {
+     labels[i] = 2
+   }
+ }
+ }
+ return(labels)
+ }
> labels = assign_labels(x, centroid1, centroid2)
> labels

```

```
[1] 1 1 1 2 2 2
```



(e) Repeat (c) and (d) until the answers obtained stop changing.

Answer->

```

> last_labels = rep(-1, 6)
> while (!all(last_labels == labels)) {
+   last_labels = labels
+   centroid1 = c(mean(x[labels==1, 1]), mean(x[labels==1, 2]))
+   centroid2 = c(mean(x[labels==2, 1]), mean(x[labels==2, 2]))
+   print(centroid1)
+   print(centroid2)
+   labels = assign_labels(x, centroid1, centroid2)
+ }

```

```
[1] 0.6666667 3.6666667
```

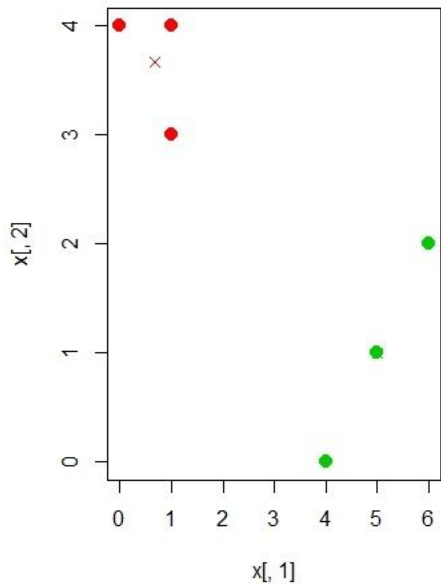
```
[1] 5 1
```

```
> labels
```

```
[1] 1 1 1 2 2 2
```

(f) In your plot from (a), color the observations according to the cluster labels obtained.

```
> plot(x[,1], x[,2], col=(labels+1), pch=20, cex=2)
> points(centroid1[1], centroid1[2], col=2, pch=4)
> points(centroid2[1], centroid2[2], col=3, pch=4)
```



9. Consider the USArrests data. We will now perform hierarchical clustering on the states.

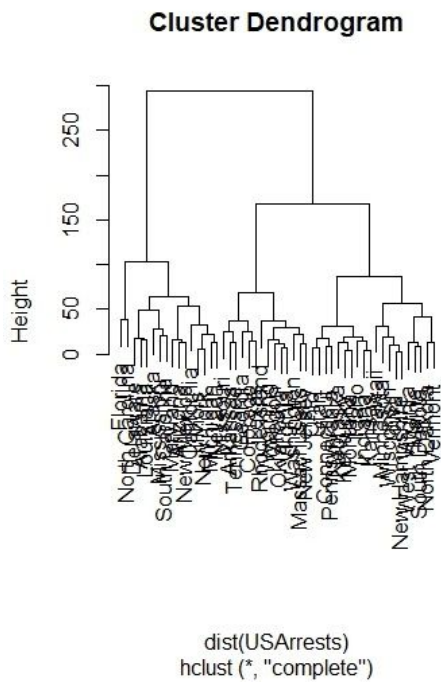
(a) Using hierarchical clustering with complete linkage and Euclidean distance, cluster the states.

Answer->

```
> set.seed(2)
```

```
> hc.complete = hclust(dist(USArrests), method="complete")
```

```
> plot(hc.complete)
```



(b) Cut the dendrogram at a height that results in three distinct clusters. Which states belong to which clusters?

Answer->

```
> cutree(hc.complete, 3)
```

Alabama	Alaska	Arizona	Arkansas
1	1	1	2
California	Colorado	Connecticut	Delaware
1	2	3	1
Florida	Georgia	Hawaii	Idaho
1	2	3	3
Illinois	Indiana	Iowa	Kansas
1	3	3	3
Kentucky	Louisiana	Maine	Maryland
3	1	3	1
Massachusetts	Michigan	Minnesota	Mississippi
2	1	3	1
Missouri	Montana	Nebraska	Nevada
2	3	3	1
New Hampshire	New Jersey	New Mexico	New York
3	2	1	1
North Carolina	North Dakota	Ohio	Oklahoma
1	3	3	2
Oregon	Pennsylvania	Rhode Island	South Carolina
2	3	2	1
South Dakota	Tennessee	Texas	Utah
3	2	2	3
Vermont	Virginia	Washington	West Virginia
3	2	2	3
Wisconsin	Wyoming		
3	2		

```
> table(cutree(hc.complete, 3))
```

```
1 2 3
```

```
16 14 20
```

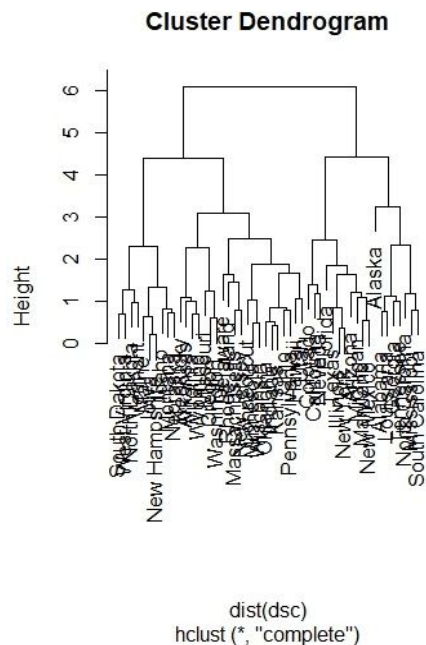
(c) Hierarchically cluster the states using complete linkage and Euclidean distance, after scaling the variables to have standard deviation one.

Answer->

```
> dsc = scale(USArrests)
```

```
> hc.s.complete = hclust(dist(dsc), method="complete")
```

```
> plot(hc.s.complete)
```



(d) What effect does scaling the variables have on the hierarchical clustering obtained? In your opinion, should the variables be scaled before the inter-observation dissimilarities are computed? Provide a justification for your answer.

Answer-> Scaling the variables effects the max height of the dendogram obtained from hierarchical clustering. From a cursory glance, it doesn't effect the bushiness of the tree obtained. However, it does affect the clusters obtained from cutting the dendogram into 3 clusters. In my opinion, for this data set the data should be standardized because the data measured has different units (UrbanPop compared to other three columns).

```
> cutree(hc.s.complete, 3)
```

Alabama	Alaska	Arizona	Arkansas
1	1	2	3
California	Colorado	Connecticut	Delaware
2	2	3	3
Florida	Georgia	Hawaii	Idaho
2	1	3	3
Illinois	Indiana	Iowa	Kansas
2	3	3	3
Kentucky	Louisiana	Maine	Maryland
3	1	3	2
Massachusetts	Michigan	Minnesota	Mississippi
3	2	3	1
Missouri	Montana	Nebraska	Nevada
3	3	3	2
New Hampshire	New Jersey	New Mexico	New York
3	3	2	2
North Carolina	North Dakota	Ohio	Oklahoma
1	3	3	3
Oregon	Pennsylvania	Rhode Island	South Carolina
3	3	3	1
South Dakota	Tennessee	Texas	Utah
3	1	2	3
Vermont	Virginia	Washington	West Virginia
3	3	3	3
Wisconsin	Wyoming		

```

      3      3
> table(cutree(hc.s.complete, 3))
 1  2  3
 8 11 31
> table(cutree(hc.s.complete, 3), cutree(hc.complete, 3))
  1  2  3
 1  6  2  0
 2  9  2  0
 3  1 10 20

```

10. In this problem, you will generate simulated data, and then perform PCA and K-means clustering on the data.

(a) Generate a simulated data set with 20 observations in each of three classes (i.e. 60 observations total), and 50 variables. Hint: There are a number of functions in R that you can use to generate data. One example is the `rnorm()` function; `runif()` is another option. Be sure to add a mean shift to the observations in each class so that there are three distinct classes.

Answer->

```

set.seed(42)
data= matrix(sapply(1:3,function(x){ rnorm(20*50,mean = 10*sqrt(x)) }),ncol=50)  # 20 obs. in
each class with 50 features.
class=unlist(lapply(1:3,function(x){rep(x,20)}))

```

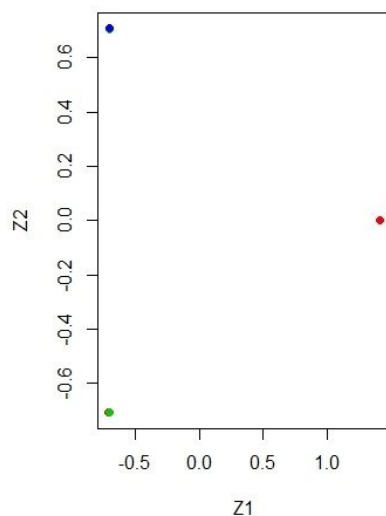
(b) Perform PCA on the 60 observations and plot the first two principal component score vectors. Use a different color to indicate the observations in each of the three classes. If the three classes appear separated in this plot, then continue on to part (c). If not, then return to part (a) and modify the simulation so that there is greater separation between the three classes. Do not continue to part (c) until the three classes show at least some separation in the first two principal component score vectors.

Answer->

```

> pca.out = prcomp(x)
> plot(pca.out$x[,1:2], col=2:4, xlab="Z1", ylab="Z2", pch=19)

```



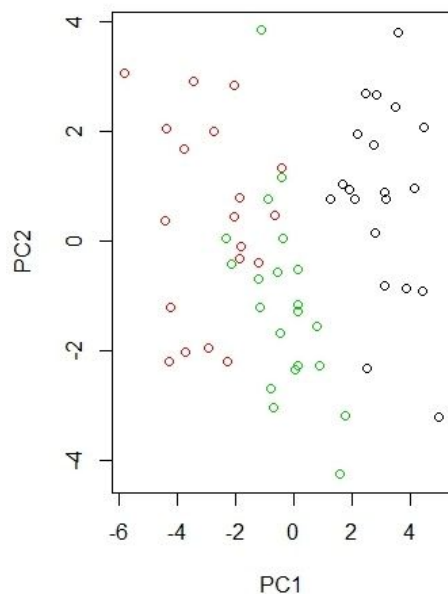
(c) Perform K-means clustering of the observations with $K = 3$. How well do the clusters that you obtained in K-means clustering compare to the true class labels? Hint: You can use the `table()`

function in R to compare the true class labels to the class labels obtained by clustering. Be careful how you interpret the results: K-means clustering will arbitrarily number the clusters, so you cannot simply check whether the true class labels and clustering labels are the same.

Answer->

```
> set.seed(1)
> kmeans.out=kmeans(data,3)
> table(kmeans.out$cluster)
 1  2  3
20 19 21
> table(class)
class
 1  2  3
20 20 20
```

We can see that there is only one observation that is miss classified.



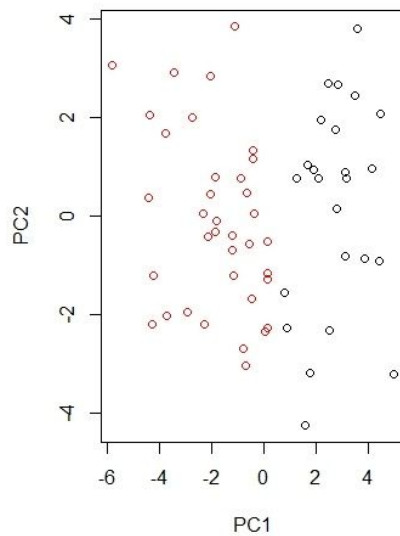
(d) Perform K-means clustering with $K = 2$. Describe your results

Answer->

```
> set.seed(1)
> kmeans.out=kmeans(data,2)
> table(kmeans.out$cluster)
 1  2
24 36
> table(class)
class
 1  2  3
20 20 20
```

K-means seem to find a single cluster that is the same as before. This can clearly be observed in the picture below as the red cluster closely matches the original green cluster.

```
> plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
```

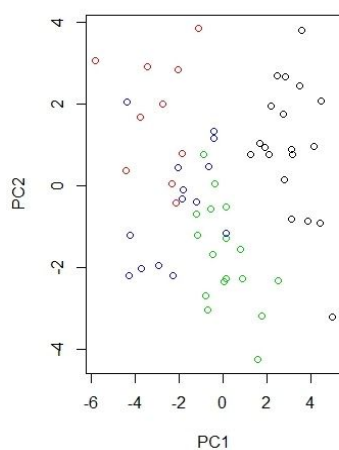
(e) Now perform K-means clustering with $K = 4$, and describe your results.

Answer->

```
> set.seed(1)
> kmeans.out=kmeans(data,4)
> table(kmeans.out$cluster)
 1  2  3  4
19 10 17 14
> table(class)
class
 1  2  3
20 20 20
```

When using 4 clusters it becomes more difficult to determine the difference between the new found clusters and the actual class values. However, by examining the plot we can see that it again find the original green cluster with some overlap between it and the remaining ones. Overlap between clusters in the two principal components is also clear, as should be expected since they may be close in the remaining dimensions.

```
> plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
```



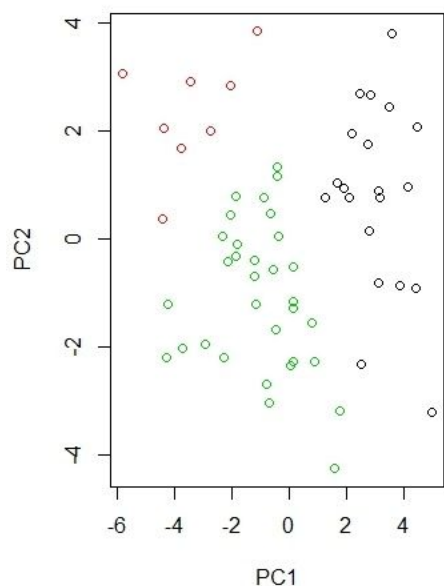
(f) Now perform K-means clustering with $K = 3$ on the first two principal component score vectors, rather than on the raw data. That is, perform K-means clustering on the 60×2 matrix of which the first column is the first principal component score vector, and the second column is the second principal component score vector. Comment on the results.

Answer->

```
> set.seed(1)
> kmeans.out=kmeans(pr.out$x[,c(1,2)],3)
> table(kmeans.out$cluster)
 1  2  3
20  8 32
> table(class)
class
 1  2  3
20 20 20
```

The algorithm performs well when clustering on the first two principal components, however, since it is missing information about the remaining dimensions observations that are close in the first two components are assigned to the same cluster which leads to mistakes. By examining the plot as before we can see that this is true, as there is no overlap.

```
> plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
```



(g) Using the scale() function, perform K-means clustering with $K = 3$ on the data after scaling each variable to have standard deviation one. How do these results compare to those obtained in (b)? Explain.

Answer->

```
> set.seed(1)
> kmeans.out=kmeans(scale(data,center = T,scale = T),3)
> table(kmeans.out$cluster)
 1  2  3
32 14 14
> table(class)
class
```

1 2 3

20 20 20

There is significant overlap in the first two clusters, and the algorithm performs poorly.

```
> plot(pr.out$x[,c(1,2)],col=kmeans.out$cluster)
```

