

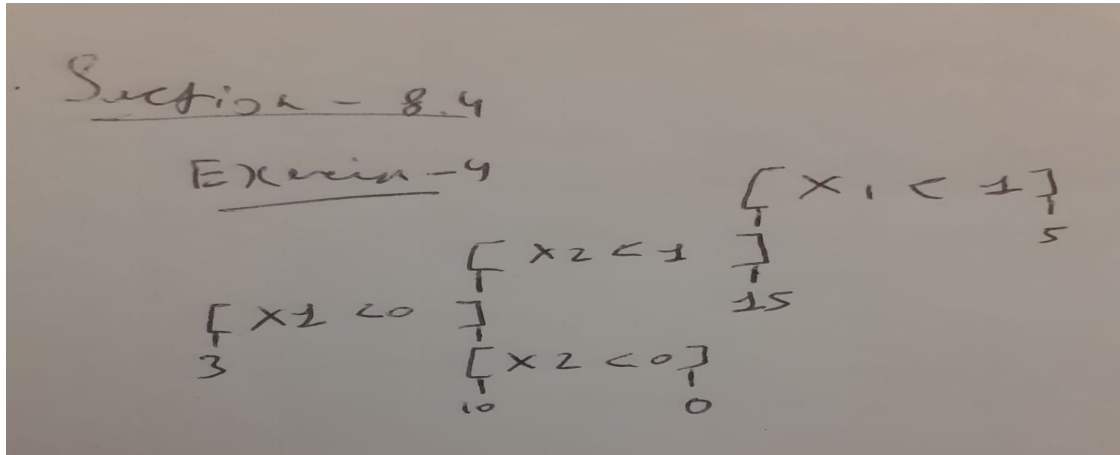
Homework-5

tg289@njit.edu

1. Exercises 4, 8, and 10 from Section 8.4 of our textbook.

Q4)


A)



B) Code-

```
> par(xpd=NA)
> plot(NA,NA,type='n',xlim=c(-2,2),ylim=c(-3,3),xlab='X1',ylab='X2')
> #X2 < 1
> lines(x=c(-2,2),y= c(1,1))
> #X1<1 with X2<1
> lines(x=c(1,1),y=c(-3,1))
> text(x=(-2+1)/2, y=-1, labels=c(-1.8))
> text(x=1.5, y=-1, labels=c(0.63))
> #X2<2 with X2>=1
> lines(x=c(-2,2),y=c(2,2))
> text(x=0,y=2.5, labels= c(2.49))
> #X1<0 with X2<2 and X2>=1
> lines(x= c(0,0),y=c(1,2))
> text(x=-1,y=1.5, labels=c(-1.06))
> text(x=1,y=1.5,labels=c(0.21))
```

Output-



Q8)

A) **code-**

```
library(ISLR)
set.seed(1)
train<- sample(1:nrow(Carseats),nrow(Carseats)/2)
carseats.train<- Carseats[train, ]
carseats.test<- Carseats[-train,
```

B) **Regression Tree->**

code-

```
library(tree)
tree.carseats<- tree(Sales~ ., data= carseats.train)
summary(tree.carseats)
```

Output-

Regression tree:

```
tree(formula = Sales ~ ., data = carseats.train)
```

Variables actually used in tree construction:

[1] "ShelveLoc" "Price"

[3] "Age" "Advertising"

[5] "CompPrice" "US"

Number of terminal nodes: 18

Residual mean deviance: 2.167 = 394.3 / 182

Distribution of residuals:

Min. 1st Qu. Median Mean

-3.88200 -0.88200 -0.08712 0.00000

3rd Qu. Max.


0.89590 4.09900

Plotting the tree

Code-

```
plot(tree.carseats)
text(tree.carseats, pretty= 0)
```

output-



Error rate-

Code-

```
> errorrate<- predict(tree.carseats, newdata= carseats.test)
> mean((errorrate- carseats.test$Sales)^2)
```

output-


[1] 4.922039

The test MSE is about **4.93**

C) **code-**

```
cv.carseats= cv.tree(tree.carseats, FUN= prune.tree)
par(mfrow= c(1,2))
plot(cv.carseats$size, cv.carseats$dev, type='b')
plot(cv.carseats$k, cv.carseats$dev, type='b')
```


Output-



Code- (best size for tree-9)

```
pruned.carseats<- prune.tree(tree.carseats, best=9)
par(mfrow= c(1,1))
plot(pruned.carseats)
text(pruned.carseats, pretty=0)
```

Output-



code-(prune tree)

```
pred.prune<- predict(pruned.carseats,carseats.test)
mean((carseats.test$Sales- pred.prune)^2)
```

Output-

[1] 5.113254

When we prune the tree the **test MSE rises to about 5.12**

D) **Code-**

```
library(randomForest)
bagging.carseats<- randomForest(Sales~ ., data = carseats.train, mtry= 10, ntree = 500, importance= TRUE)
bagging.pred<- predict(bagging.carseats,carseats.test)
mean((carseats.test$Sales- bagging.pred)^2)
```

Output-

```
[1] 2.622982
```

We can see that when we use bagging the test **MSE decreases to 2.62**

Code-(for importance)-

```
importance(bagging.carseats)
```

Output-

	%IncMSE	IncNodePurity
CompPrice	25.8083065	176.750788
Income	4.3976727	90.755860
Advertising	12.2530540	96.396343
Population	-3.1644545	55.382276
Price	56.2028471	501.877790
ShelveLoc	46.4048838	377.701755
Age	16.0014582	155.431845
Education	-0.2585088	44.103966
Urban	-0.3479058	8.679922
US	5.5939322	18.643493

We can conclude from this - **Price, ShelveLoc and CompPrice are three most important variables**

E) **Code-**

```
random_f<- randomForest(Sales ~., data= carseats.train, mtry=3,ntree=500,importance=TRUE)
random_pred<-predict(random_f,carseats.test)
mean((carseats.test$Sales- random_pred)^2)
```

Output-

```
[1] 3.070883
```

Code-

```
importance(random_f)
```

Output-

	%IncMSE	IncNodePurity
CompPrice	15.3620526	156.79771
Income	3.1692502	120.90880
Advertising	7.8046269	108.57608
Population	-2.5512335	100.56722
Price	36.5947531	396.36763
ShelveLoc	35.9522295	290.53078
Age	11.4111604	174.34911
Education	0.4647345	72.66570
Urban	1.4222554	15.61579
US	6.3063458	34.71705

We again see that price shevlock and compprice are the most important variables

Q10)

A) **Code-**

```
library(ISLR)
Hitters<- na.omit(Hitters)
Hitters$Salary<- log(Hitters$Salary)
```


B) Code-

```
train<- 1:200  
Hitters.train= Hitters[train,]  
Hitters.test= Hitters[-train,]
```

C) Code-

```
install.packages('gbm')  
library(gbm)  
set.seed(103)  
pows=seq(-10, -0.2, by=0.1)  
lambdas= 10^pows  
length.lambdas= length(lambdas)  
train.errors= rep(NA, length.lambdas)  
test.errors= rep(NA,length.lambdas)  
for(i in 1:length.lambdas){  
  boost.hitters= gbm(Salary~ ., data= Hitters.train, distribution = 'gaussian',  
    n.trees= 1000, shrinkage= lambdas[i])  
  train.pred= predict(boost.hitters, Hitters.train, n.trees= 1000)  
  test.pred= predict(boost.hitters,Hitters.test, n.trees= 1000)  
  train.errors[i]= mean((Hitters.train$Salary- train.pred)^2)  
  test.errors[i]= mean((Hitters.test$Salary- test.pred)^2)  
}  
plot(lambdas, train.errors, type='b',xlab='Shrinkage',ylab='Train MSE',col='blue',pch=20)
```


Output-



D) Code-

```
plot(lambdas,test.errors, type='b',xlab='Shrinkage',ylab='Test MSE',col='blue',pch=20 )
```

Output-



Code- (For MSE)

```
min(test.errors)
```

Output-

```
[1] 0.008058291
```

Code-(For lambda)

```
lambdas[which.min(test.errors)]
```

Output-

```
[1] 0.05011872
```

The minimum test error is 0.008 and corresponding lambda value is 0.05

E) Code-

```
library(glmnet)
lm.fit= lm(Salary~ ., data= Hitters.train)
lm.pred= predict(lm.fit, Hitters.test)
mean((Hitters.test$Salary- lm.pred)^2)
```

Output-

```
[1] 0.01496996
```

Code-

```
set.seed(1)
x= model.matrix(Salary ~., data= Hitters.train)
y= Hitters.train$Salary
x.test= model.matrix(Salary~ ., data= Hitters.test)
lasso.fit= glmnet(x,y,alpha=1)
lasso.pred= predict(lasso.fit, s=0.01, newx= x.test)
mean((Hitters.test$Salary- lasso.pred)^2)
```

Output-

```
[1] 0.01377199
```

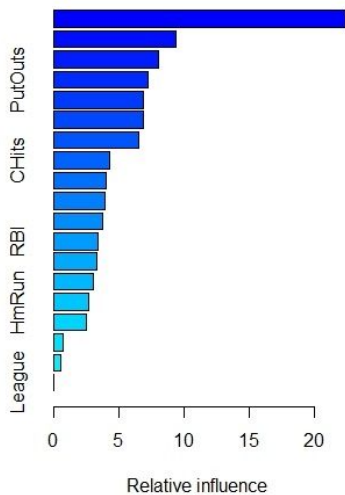
Both linear model and regularization like Lasso have higher test MSE than boosting

F) Code-

```
library(gbm)
boosting.best<- gbm(Salary~ ., data= Hitters.train, distribution = 'gaussian',
n.trees= 1000, shrinkage= lambdas[which.min(test.errors)])
summary(boosting.best)
```

Output-

	var	rel.inf
CAtBat	CAtBat	20.1133386
CRuns	CRuns	12.4546958
CRBI	CRBI	10.9286391
PutOuts	PutOuts	7.0217036
CWalks	CWalks	6.8343354
Years	Years	6.3477290
Walks	Walks	5.7228669
CHmRun	CHmRun	5.4150877
RBI	RBI	4.6494594
Hits	Hits	4.4069403
Assists	Assists	3.4818646
AtBat	AtBat	2.9389697
HmRun	HmRun	2.6833097
CHits	CHits	2.6574682
Errors	Errors	1.6588080
Runs	Runs	1.2466717
NewLeague	NewLeague	0.6987026
Division	Division	0.5333779
League	League	0.2060317

**G) Code-**

```
library(randomForest)
set.seed(25)
rf.hitters<- randomForest(Salary~ ., data= Hitters.train, ntree=500, mtry=19)
rf.pred<- predict(rf.hitters, Hitters.test)
mean((Hitters.test$Salary- rf.pred)^2)
```

Output-

```
[1] 0.007193393
```


2. Exercises 3 and 7 from Section 9.7 of our textbook.

Q3)

A) Code-

```
x1= c(3,2,4,1,2,4,4)
x2= c(4,2,4,4,1,3,1)
colors=c('blue','blue','blue','blue','red','red','red'))
plot(x1,x2,col=colors,xlim=c(0,6),ylim=c(0,6))
```

Output-



B) Equation-

Section - 9.7

Exercise - 3 (5)

The horizontal margin classifier has to be in 6/w obs. #2, #3 are #5 #6

$$(2, 2), (4, 4)$$

$$(2, 1), (4, 3)$$

$$\bar{x} = \left(\frac{2+4}{2}, \frac{2+4}{2} \right), \left(\frac{4+2}{2}, \frac{4+3}{2} \right)$$

$$= (2, 1.5), (4, 3.5)$$

$$b = (3.5 - 1.5) / (4 - 2) = \boxed{1}$$


$$a = x_2 - x_1 = 1.1 - 2 = \boxed{-0.5}$$

Sketch-

Code-

```
plot(x1,x2,col=colors,xlim=c(0,5),ylim=c(0,5))  
abline(-0.5,1)
```

Output-




C) The classification rules are - $0.5 - x_1 + x_2 > 0$

D) Code-

```
plot(x1,x2,col=colors,xlim=c(0,5),ylim=c(0,5))  
abline(-0.5,1)  
abline(-1,1,lty=2)  
abline(0,1,lty=2)
```


Output-



E) **Code-**

```
plot(x1,x2,col=colors,xlim=c(0,5),ylim=c(0,5))
abline(-0.5,1)
arrows(2,1,2,1.5)
arrows(2,2,2,1.5)
arrows(4,4,4,3.5)
arrows(4,3,4,3.5)
```

Output-




F) A slight movement of observation #7 (4,1) red would not have an effect on the maximal margin hyperplane since its movement would be outside of the margin.

G) **Code-**

```
plot(x1,x2,col=colors,xlim=c(0,5),ylim=c(0,5))
abline(-.8,1)
```

Output-




Equation of the hyperplane- $(-0.8 \cdot X_1 + X_2 > 0)$

H) Code-

```
plot(x1,x2,col=colors,xlim=c(0,5),ylim=c(0,5))
points(c(4),c(2),col=c("blue"))
```

Output-



Q7)

A) Code-

```
library(ISLR)
gas.median= (median(Auto$mpg))
factor_med= ifelse(Auto$mpg> gas.median, 1, 0)
Auto$mpglevel= as.factor(factor_med)
```

B) Code-

```
library(e1071)
set.seed(3250)
tune.in<- tune(svm,mpglevel~ ., data=Auto, kernel='linear', ranges= list(cost= c(0.01,0.1,1,5,10,100)))
summary(tune.in)
```

Output-

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost

1

- best performance: 0.01025641

- Detailed performance results:

	cost	error	dispersion
1	1e-02	0.07410256	0.03724108
2	1e-01	0.04596154	0.02359743
3	1e+00	0.01025641	0.01792836
4	5e+00	0.01782051	0.02088734
5	1e+01	0.02038462	0.02001214
6	1e+02	0.03307692	0.02397013

We see that cross-validation error is minimized for **cost=1**.

C) Code-

```
set.seed(22)
tune.out= tune(svm, mpglevel~ ., data=Auto, kernel= "polynomial", ranges=
list(cost=c(0.1,1,5,10),degree=c(2,3,4)))
summary(tune.out)
```

Output-

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost degree

10 2

- best performance: 0.5510897

- Detailed performance results:

	cost	degree	error	dispersion
1	0.1	2	0.5741026	0.02415229
2	1.0	2	0.5741026	0.02415229
3	5.0	2	0.5741026	0.02415229
4	10.0	2	0.5510897	0.04591189
5	0.1	3	0.5741026	0.02415229
6	1.0	3	0.5741026	0.02415229
7	5.0	3	0.5741026	0.02415229
8	10.0	3	0.5741026	0.02415229
9	0.1	4	0.5741026	0.02415229
10	1.0	4	0.5741026	0.02415229
11	5.0	4	0.5741026	0.02415229
12	10.0	4	0.5741026	0.02415229

#The lowest cross-validation error is obtained for cost=10 and degree=2.

Code-

```
set.seed(461)
```

```
tune.out<- tune(svm, mpglevel~ ., data=Auto, kernel='radial',ranges=list(cost= c(0.1,1,5,10), gamma= c(0.01,.,1,1,5,10,100)))
```

```
summary(tune.out)
```

Output-

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

cost gamma

10 0.1

- best performance: 0.02551282

- Detailed performance results:

	cost	gamma	error	dispersion
1	0.1	1e-02	0.08666667	0.04967643
2	1.0	1e-02	0.07391026	0.04561258
3	5.0	1e-02	0.04839744	0.03496088
4	10.0	1e-02	0.03064103	0.02651089
5	0.1	1e-01	0.07903846	0.04876296
6	1.0	1e-01	0.05346154	0.04057827
7	5.0	1e-01	0.02807692	0.02551393
8	10.0	1e-01	0.02551282	0.02702937
9	0.1	1e+00	0.57416667	0.05205090
10	1.0	1e+00	0.06352564	0.04307392
11	5.0	1e+00	0.06358974	0.03814303
12	10.0	1e+00	0.06608974	0.04146890
13	0.1	5e+00	0.57416667	0.05205090
14	1.0	5e+00	0.52814103	0.05853644
15	5.0	5e+00	0.51538462	0.05940094
16	10.0	5e+00	0.51538462	0.05940094
17	0.1	1e+01	0.57416667	0.05205090

```

18 1.0 1e+01 0.53346154 0.06755706
19 5.0 1e+01 0.53089744 0.06238863
20 10.0 1e+01 0.53089744 0.06238863
21 0.1 1e+02 0.57416667 0.05205090
22 1.0 1e+02 0.57416667 0.05205090
23 5.0 1e+02 0.57416667 0.05205090
24 10.0 1e+02 0.57416667 0.05205090

```

#Finally, for radial basis kernel, cost=10 and gamma=0.01.


D) Code-(for linear SVM)

```

svm.linear<- svm(mpglevel~ ., data=Auto,kernel="linear",cost=1)
svm.poly= svm(mpglevel~ ., data= Auto, kernel='polynomial',cost=10,degree=2)
svm.radical= svm(mpglevel~ ., data=Auto, kernel='radial',cost=10,gamma=0.01)
plotpairs<- function(fit){
  for(name in names(Auto)[!(names(Auto) %in% c('mpg','mpglevel','name'))]) {
    plot(fit,Auto,as.formula(paste('mpg~',name,sep="")))
  }
}
plotpairs(svm.linear)

```

Output-



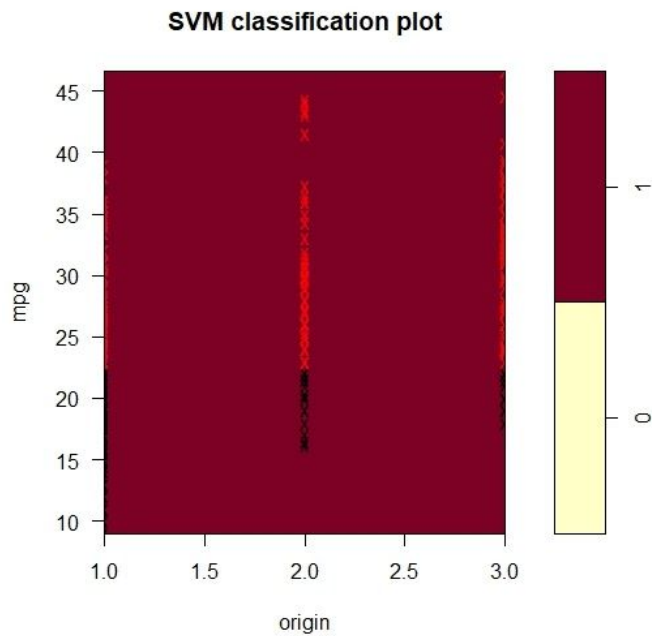
Code- (for polynomial basis)

```

svm.linear<- svm(mpglevel~ ., data=Auto,kernel="linear",cost=1)
svm.poly= svm(mpglevel~ ., data= Auto, kernel='polynomial',cost=10,degree=2)
svm.radical= svm(mpglevel~ ., data=Auto, kernel='radial',cost=10,gamma=0.01)
plotpairs<- function(fit){
  for(name in names(Auto)[!(names(Auto) %in% c('mpg','mpglevel','name'))]) {
    plot(fit,Auto,as.formula(paste('mpg~',name,sep="")))
  }
}
plotpairs(svm.poly)

```

Output-



Code-(for radial basis)

```
svm.linear<- svm(mpglevel~ ., data=Auto,kernel="linear",cost=1)
svm.poly= svm(mpglevel~ ., data= Auto, kernel='polynomial',cost=10,degree=2)
svm.radical= svm(mpglevel~ ., data=Auto, kernel='radial',cost=10,gamma=0.01)
plotpairs<- function(fit){
  for(name in names(Auto)[!(names(Auto) %in% c('mpg','mpglevel','name'))]) {
    plot(fit,Auto,as.formula(paste('mpg~',name,sep="")))
  }
}
```

plotpairs(svm.radial)

Output-

